

Evolutionary Computing

Chapter 3 – What is an Evolutionary Algorithm?

Common model of evolutionary processes

- Population of individuals
- Individuals have a fitness
- Reproduction / variation operators
 - mutation
 - recombination (a.k.a. crossover)
- Selection towards higher fitness
 - “survival of the fittest” and
 - “mating of the fittest”
- The fitness of the population increases over time

Two pillars of evolution

There are two competing forces

Increasing population
diversity by variation

- mutation
- recombination

Push towards **novelty**

Decreasing population
diversity by selection

- of parents
- of survivors

Push towards **quality**

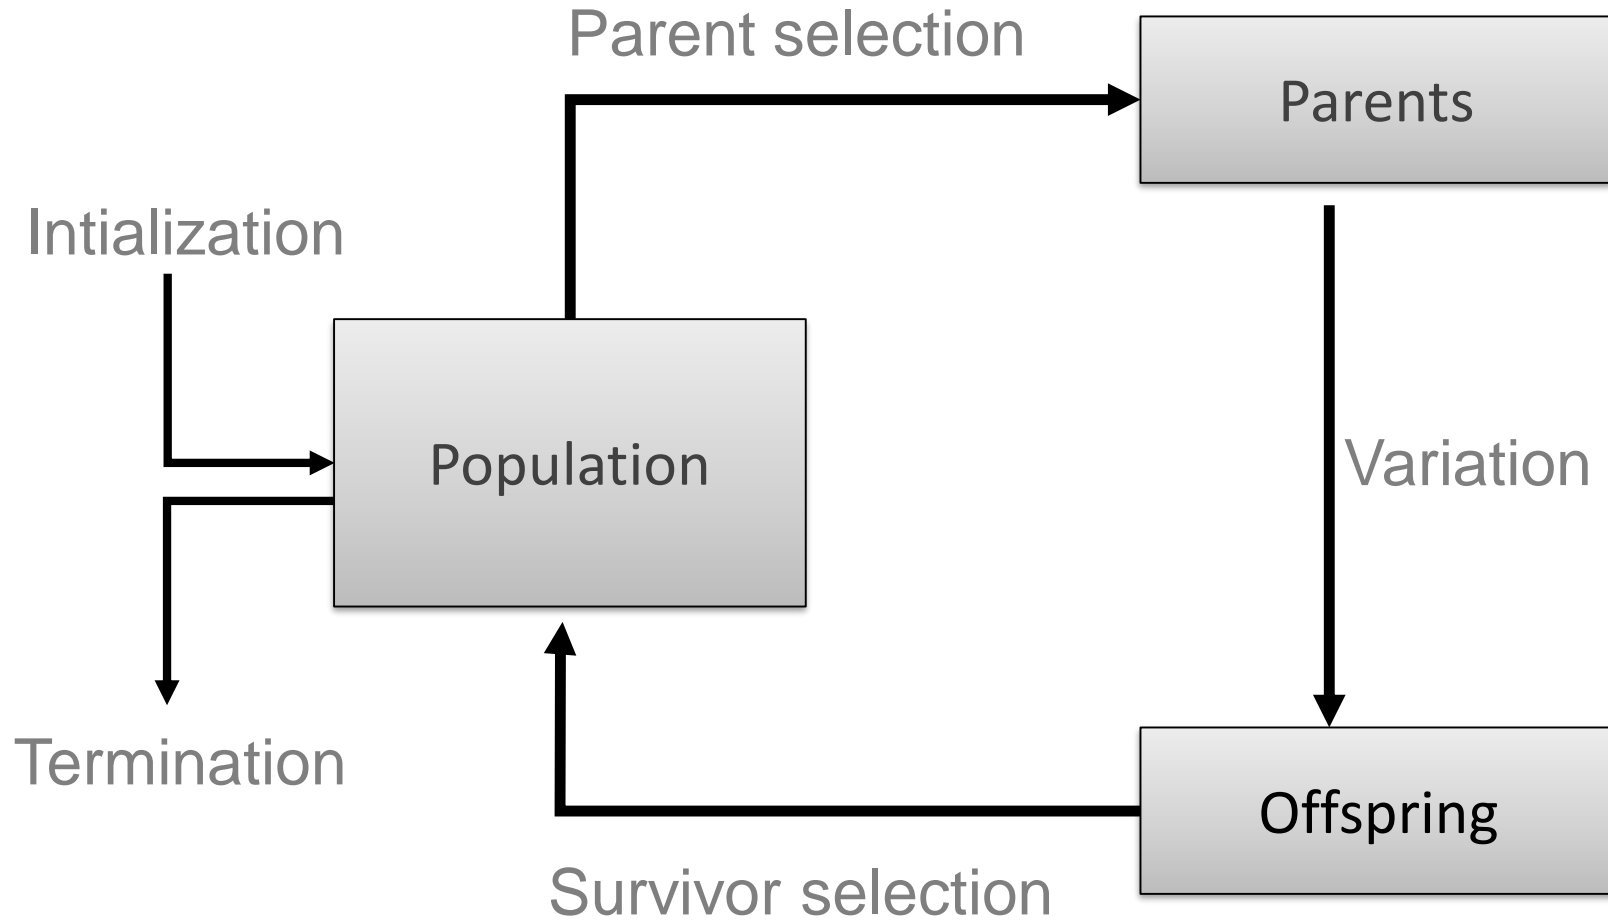
- **Selection** operators act on **population** level
- **Variation** operators act on **individual** level

What is an Evolutionary Algorithm?

- Scheme of an EA
- Main EA components:
 - Representation / evaluation / population
 - Parent selection / survivor selection
 - Mutation / recombination
- Exercise: eight-queens problem
- EA behaviour
- EAs in context

General scheme of EAs

arrows are the operations



EA scheme in pseudo-code

BEGIN

INITIALISE population with random candidate solutions;

EVALUATE each candidate;

REPEAT UNTIL (*TERMINATION CONDITION* is satisfied) DO

1 *SELECT* parents;

2 *RECOMBINE* pairs of parents;

3 *MUTATE* the resulting offspring;

4 *EVALUATE* new candidates;

5 *SELECT* individuals for the next generation;

OD

END

Main EA components: Representation

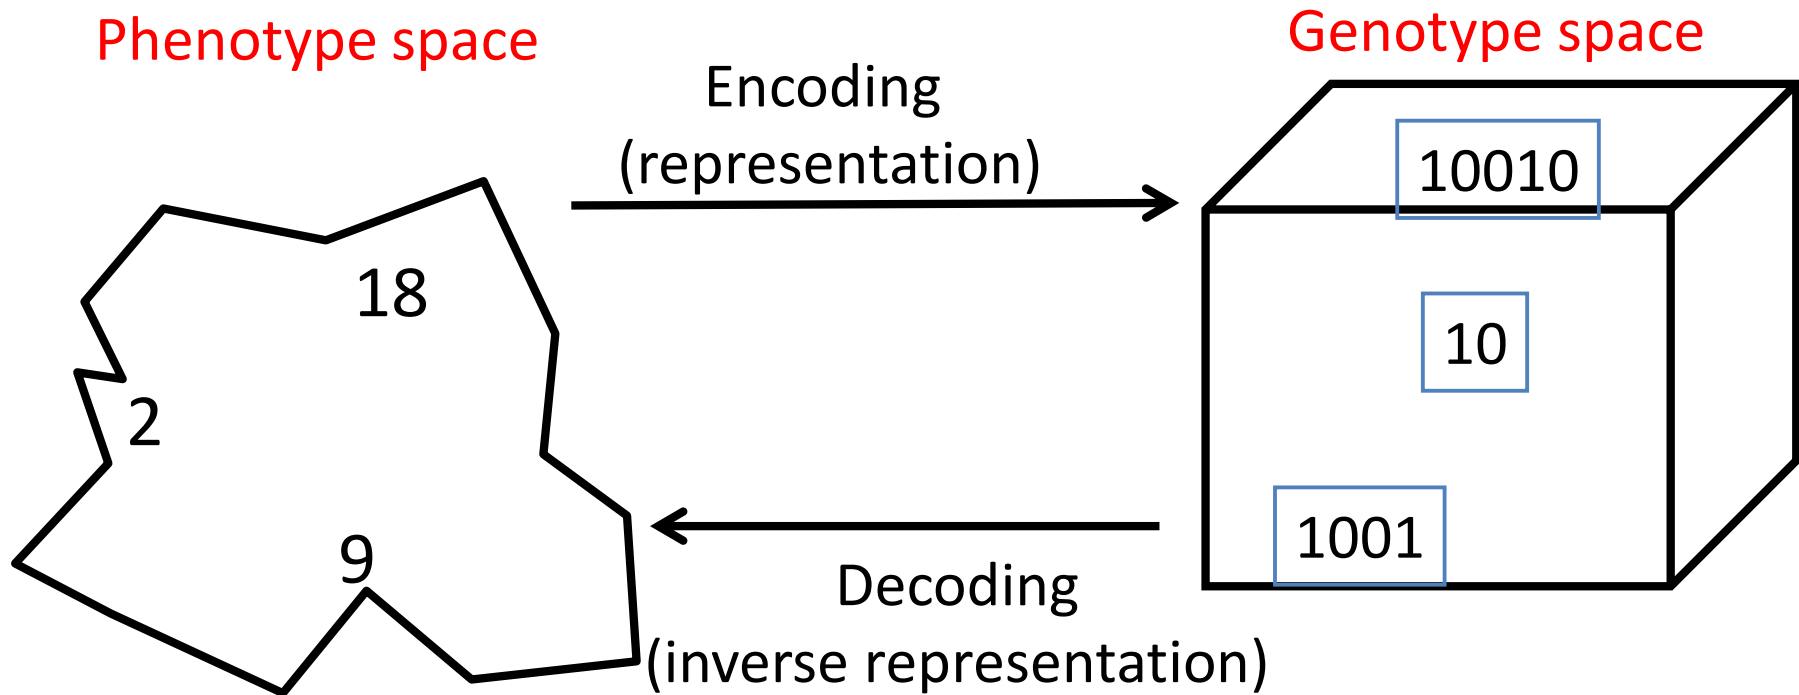
- Role: provides code for candidate solutions that can be manipulated by variation operators
- Leads to two levels of existence
 - **phenotype: object** in original problem context, the outside
 - **genotype: code** to denote that object, the inside (chromosome, “digital DNA”)
- Implies two mappings:
 - Encoding : phenotype \rightarrow genotype (not necessarily one to one)
 - Decoding : genotype \rightarrow phenotype (must be one to one)
- Chromosomes contain genes, which are in (usually fixed) positions called loci (sing. locus) and have a value (allele)

Some terminology

- Variable vs. value
- Gene vs. allele
- Example:
 - x vs. 3.14
 - color vs. red

Main EA components: Representation

Example: represent integer values by their binary code



In order to find the global optimum, every feasible solution must be represented in genotype space

Main EA components: Evaluation / fitness function

Role:

- Represents the task to solve, the requirements to adapt to (can be seen as “the environment”)
- Enables selection (provides basis for comparison)
- e.g., some phenotypic traits are advantageous, desirable, e.g. big ears cool better, these traits are rewarded by more offspring that will expectedly carry the same trait
- A.k.a. *quality* function or *objective* function
- Assigns a single real-valued fitness to each phenotype which forms the basis for selection
 - So the more discrimination (different values) the better
- Typically we talk about fitness being maximised
 - Some problems may be best posed as minimisation problems, but conversion is trivial

Main EA components: Population

- Role: holds the candidate solutions of the problem as individuals (genotypes)
- Formally, a population is a multiset of individuals, i.e. the same element might occur multiple times
- **Population is the basic unit of evolution, i.e., the population is evolving, not the individuals**
- Selection operators act on population level
- Variation operators act on individual level

Main EA components: Population

- Some sophisticated EAs also assert a spatial structure on the population e.g., a grid
- Selection operators usually take whole population into account i.e., reproductive probabilities are *relative* to current generation
- **Diversity** of a population refers to the number of different fitness values / phenotypes / genotypes present (note: not the same thing)

Main EA components: Selection

Role:

- Identifies individuals
 - to become parents
 - to survive
- Pushes population towards higher fitness
- Usually probabilistic / stochastic
 - high quality solutions more likely to be selected than low quality
 - but not guaranteed
 - even worst in current population usually has non-zero probability of being selected
- This *stochastic* nature can help escape from local optima


Main EA components: Selection

Example: roulette wheel selection

$\text{fitness}(A) = 3$

$\text{fitness}(B) = 1$

$\text{fitness}(C) = 2$



In principle, any selection mechanism can be used for parent selection as well as for survivor selection

Main EA components: Selection

- Survivor selection a.k.a. ***replacement***
- Most EAs use fixed population size so need a way of going from parents + offspring to next generation
- Often deterministic (while parent selection is usually stochastic)
 - Fitness based : e.g., rank parents + offspring and take best
 - Age based: make as many offspring as parents and delete all parents
- Sometimes a combination of stochastic and deterministic, e.g., elitism (the best n individuals always survive) is a deterministic rule

Main EA components: Variation operators

- Role: to generate new candidate solutions
- Usually divided into two types according to their arity (number of inputs):
 - Arity 1 : mutation operators
 - Arity >1 : recombination operators
 - Arity = 2 typically called crossover
 - Arity > 2 multi-parent reproduction, is possible, seldom used in EC
- There has been much debate about relative importance of recombination and mutation
 - Nowadays most EAs use both – pragmatic attitude is advisable
 - **Variation operators must match the given representation**

Main EA components: Mutation

- Role: causes small, random variations
- Acts on one genotype and delivers another
- Element of randomness is essential and differentiates it from other heuristic operators
- Importance ascribed depends on representation and historical dialect:
 - Binary GAs – background operator responsible for preserving and introducing diversity
 - EP – only search operator
 - GP – hardly used
- May guarantee connectedness of search space and hence convergence proofs (see later in Theory chapter)


Main EA components: Mutation

GENOTYPE

PHENOTYPE


before

1 1 1 1 1 1 1



after

1 1 1 0 1 1 1



Main EA components: Recombination


- Role: merges information from parents into offspring
- Choice of what information to merge is stochastic
- Most offspring may be worse, or the same as the parents
- Hope is that some are better by combining elements of genotypes that lead to good traits
- Principle has been used for millennia by breeders of plants and livestock

Main EA components: Recombination


GENOTYPE

PHENOTYPE

Parents



Offspring



Main EA components: Initialisation / Termination


- Initialisation usually done at random,
 - Need to ensure even spread and mixture of possible alleles (values)
 - Can include existing solutions, or use problem-specific heuristics, to “seed” the population
- Termination condition checked every generation
 - Reaching some (known/hoped for) fitness level
 - Reaching some maximum number of generations
 - Reaching some minimum level of diversity
 - Reaching some specified number of generations without fitness improvement (stagnation)

Different types of EAs

- Historically different flavours of EAs have been associated with different data types to represent solutions
 - Binary strings : Genetic Algorithms
 - Real-valued vectors : Evolution Strategies
 - Finite State Machines: Evolutionary Programming
 - LISP trees: Genetic Programming
- These historical differences are largely irrelevant, best strategy
 - choose representation to suit problem
 - choose variation operators to suit representation
- Selection operators only use fitness and so are independent of representation



Exercise:
solving the 8-queens problem

The 8-queens problem




Place 8 queens on an 8x8 chessboard in such a way that they cannot check each other.

The 8-queens problem: Representation



The 8-queens problem: Representation



- Table of 8 x 8
- Bit-string of length 64


But:

- More than 8 queens is an infeasible solution
- You already know that there can't be more than two queens in one row or column

The 8-queens problem: Representation

Phenotype:
a board configuration


Genotype:
a permutation of
the numbers 1–8



Possible mapping

1	3	5	2	6	4	7	8
---	---	---	---	---	---	---	---

The 8-queens problem: Representation



1	3	5	2	6	4	7	8
---	---	---	---	---	---	---	---

- Reduction of the dimension of the search space from 64 to 8
- Removal of many poor solutions

The 8-queens problem: Representation

Different variation operators are suitable for different representations.

The 8-queens problem: Fitness evaluation

- **Penalty** of one queen: the number of queens she can check
- Penalty of a configuration: the sum of penalties of all queens
- Note: penalty is to be minimized
- **Fitness** of a configuration: inverse penalty to be maximized

The 8-queens problem: Mutation


Small variation in one permutation, e.g., swapping values of two randomly chosen positions.



The 8-queens problem: Recombination

Combining two permutations into two new permutations:

- choose random crossover point
- copy first parts into children
- create second part by inserting values from other parent:
 - in the order they appear there
 - beginning after crossover point
 - skipping values already in child



The 8-queens problem: Selection

- Parent selection:
 - Pick 5 parents randomly and take best two to undergo crossover
- Survivor selection (replacement)
 - When inserting a new child into the population, choose an existing member to replace, e.g,
 - the worst individual
 - the best individual that is worse than this child:
 - sorting the whole population by decreasing fitness
 - enumerating this list from high to low
 - replacing the first with a fitness lower than the given child

The 8-queens problem: Summary

Representation	Permutations
Recombination	“Cut-and-crossfill” crossover
Recombination probability	100%
Mutation	Swap
Mutation probability	80%
Parent selection	Best 2 out of random 5
Survival selection	Replace worst
Population size	100
Number of Offspring	2
Initialisation	Random
Termination condition	Solution or 10,000 fitness evaluation

Note that this is ***only one possible*** set of choices of operators and parameter values.

Typical EA behaviour: Stages


Stages in optimising on a 1-dimensional fitness landscape



Early stage:
quasi-random population distribution




Mid-stage:
population arranged around/on hills



Late stage:
population concentrated on high hills


Working of an EA demo

Searching a fitness landscape without “niching”



Working of an EA demo

Searching a fitness landscape with “niching”




Typical EA behaviour: progression of fitness



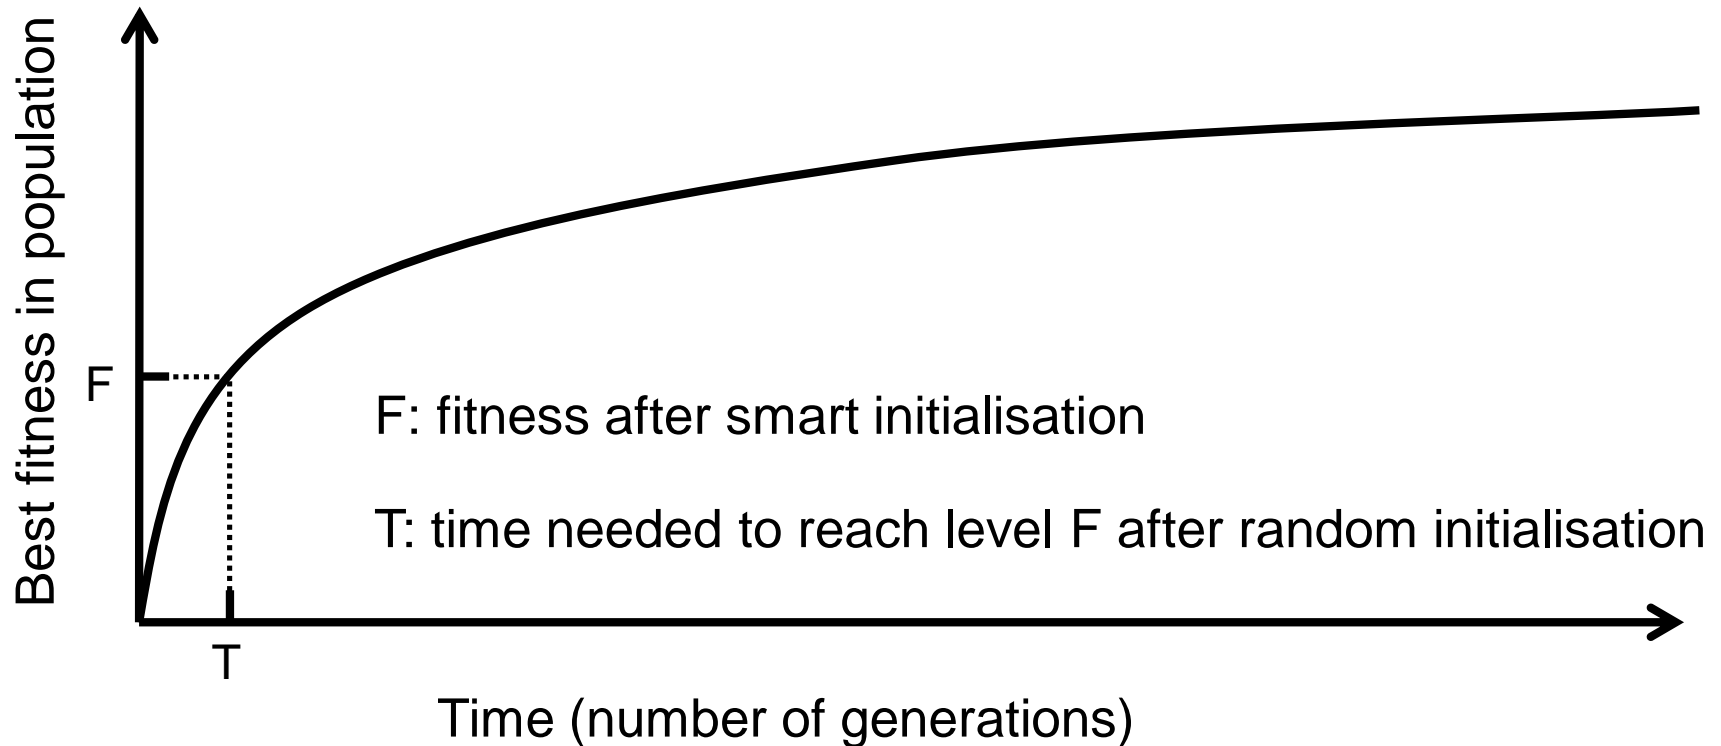
Typical run of an EA shows so-called “anytime behavior”

Are long runs beneficial?



- Answer:
 - It depends on how much you want the last bit of progress
 - May be better to do more short runs

Is it worth expending effort on smart initialisation?




- Answer: it depends.
 - Possibly good, if good solutions/methods exist.
 - May not be necessary, EA will catch up quickly
 - Care is needed, see chapter/lecture on hybridisation.

Evolutionary Algorithms in context

- EAs fall into the category of generate-and-test algorithms, a.k.a. trial-and-error algorithms
- They are stochastic, population-based search algorithms
- For most problems a problem-specific tool may:
 - perform better than a generic search algorithm on most instances,
 - have limited utility,
 - not do well on all instances
- EAs are meant to provide robust tool that shows:
 - evenly good performance
 - over a range of problems and instances


EAs as problem solvers: Goldberg view (1989)



EAs and domain knowledge

- Trend in the 90's: adding problem specific knowledge to EAs
 - (special variation operators, repair, etc)
- Result: EA performance curve “deformation”:
 - better on problems of the given type
 - worse on problems different from given type
 - amount of added knowledge is variable
- Theory suggests the search for an “all-purpose” algorithm may be fruitless

EAs as problem solvers: Michalewicz view (1996)



EC and global optimisation

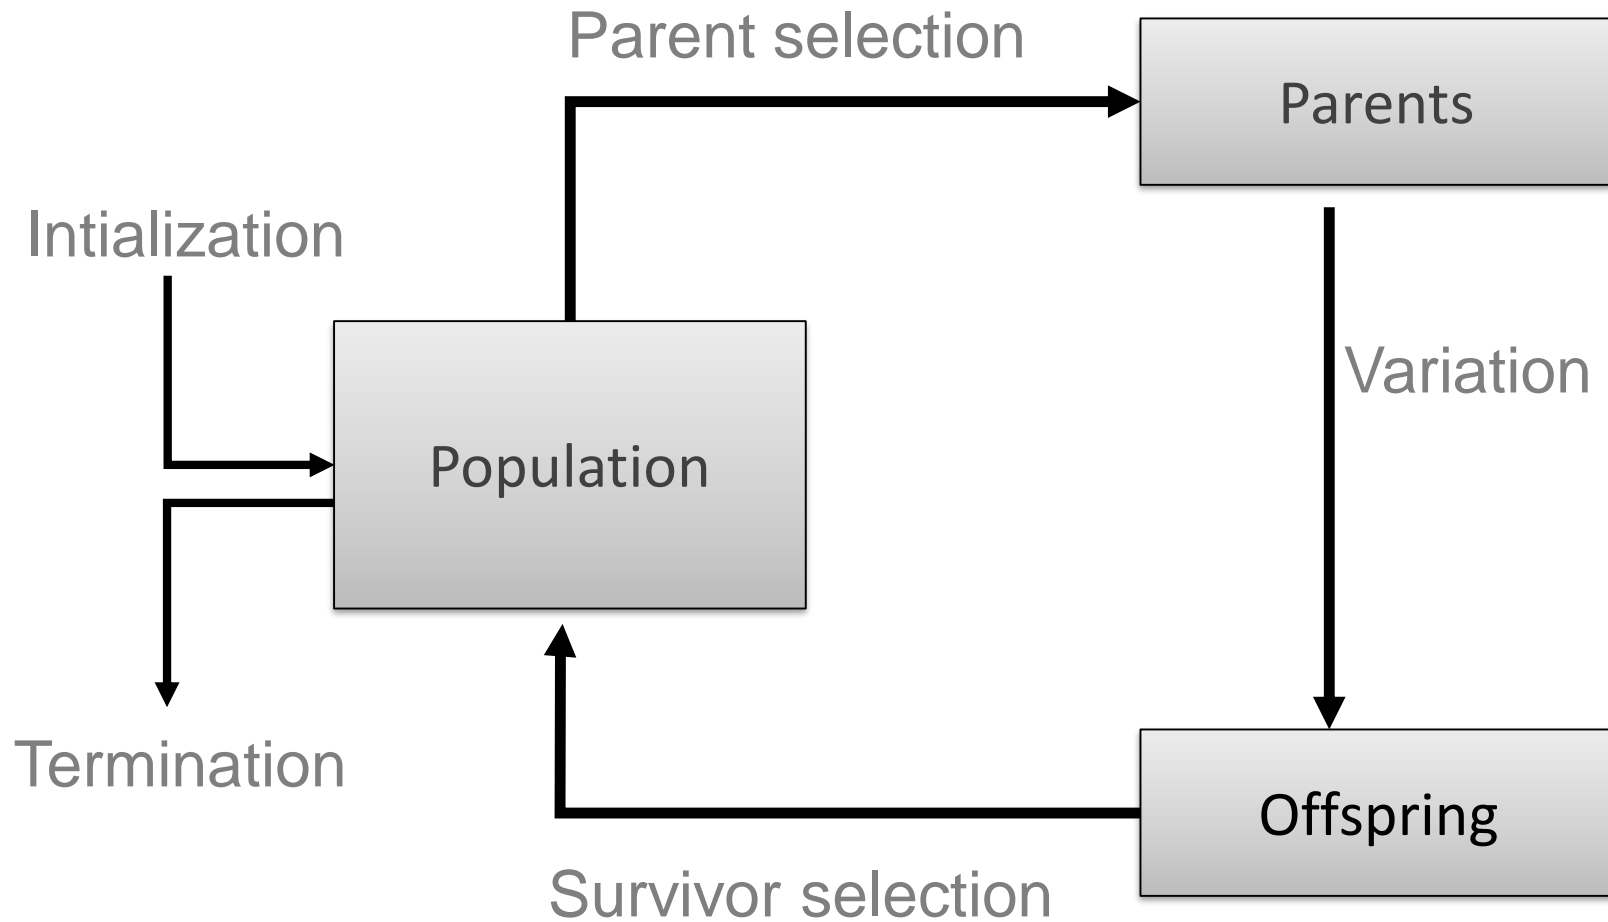
- Global optimisation: search for finding best solution x^* out of some fixed set S
- Deterministic approaches
 - e.g. box decomposition (branch and bound etc)
 - Guarantee to find x^* ,
 - May have bounds on runtime, usually super-polynomial
- Heuristic approaches (generate and test)
 - rules for deciding which $x \in S$ to generate next
 - no guarantees that best solutions found are globally optimal
 - no bounds on runtime

EC and neighbourhood search

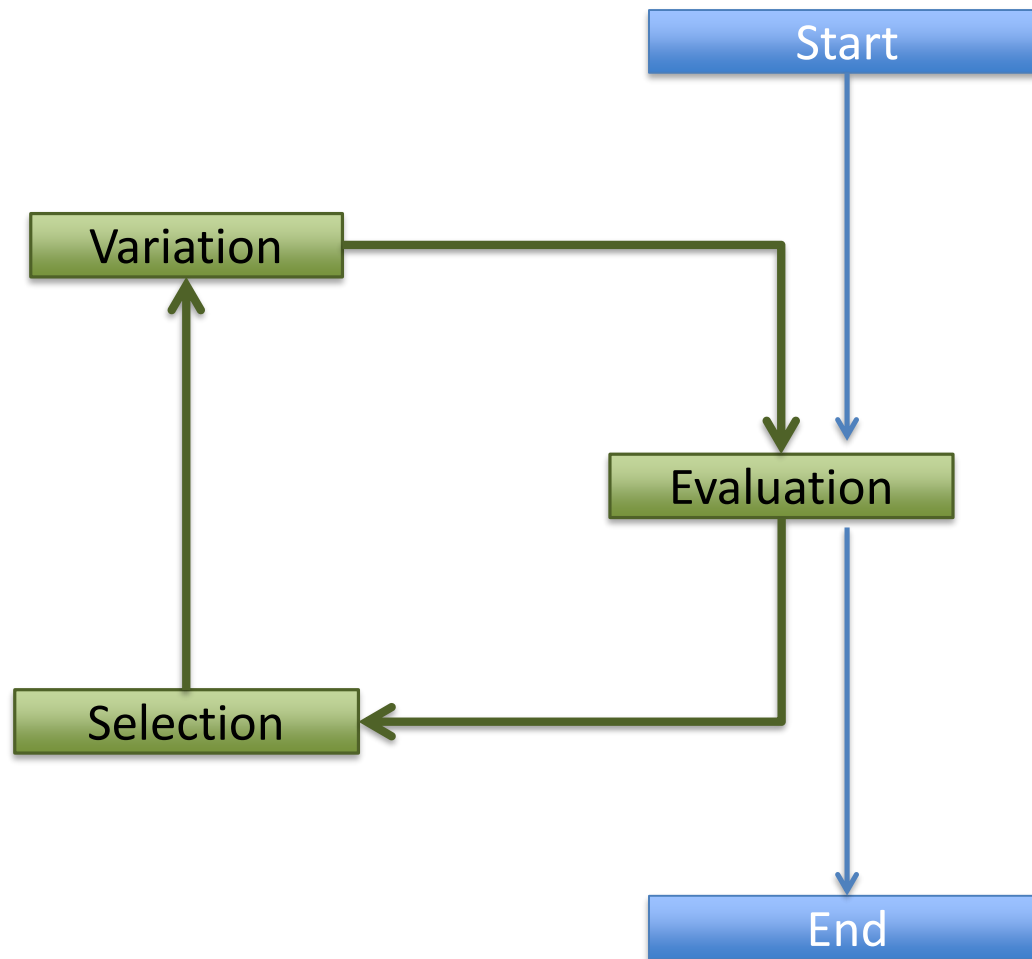
- Many heuristics impose a neighbourhood structure on S
- Such heuristics may guarantee that best point found is *locally optimal* e.g. Hill-Climbers:
 - But problems often exhibit many local optima
 - Often very quick to identify good solutions
- EAs are distinguished by (the combination of) the following:
 - Use of population
 - Use of multiple, stochastic search operators
 - Variation operators with arity >1 , thus combining information of more candidate solutions
 - Stochastic selection

Question: what is a neighbourhood in an EA?


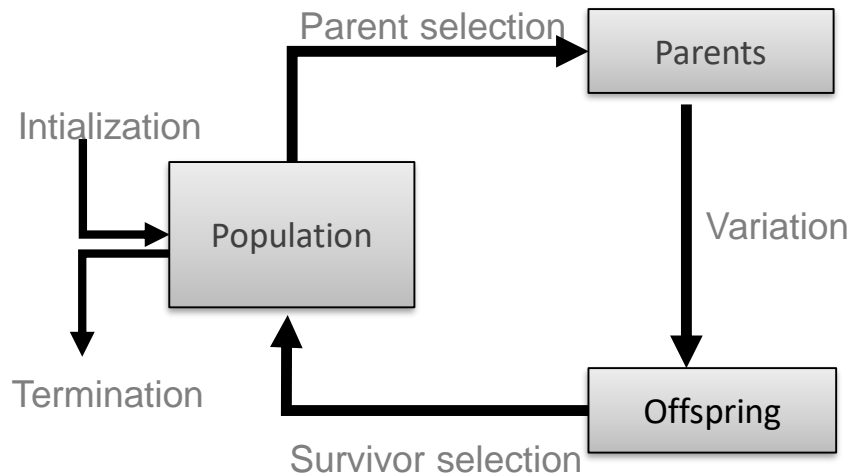
General scheme of EAs – version 1



General scheme of EAs – version 2



Q: What is the difference?



Important points:

- There are different diagrams to represent an EA scheme
- Variation operators push towards novelty and selection operators push towards quality
- Variation operators need to match the representation, selection operators are independent from the representation (hence, from the problem at hand)
- Selection operators act on the population level and variation operators on the individual level